

docker

Docker 101

Agenda

- About you & me
- What is Docker
- Why use Docker
- Writing Dockerfiles
- Running Containers
- **BREAK**
- Tools
- Resources
- Ways to integrate PHP with Docker
- Conclusion & Recap

About you

- What is your job?
- In which programming language are
 - you most comfortable in?
 - not comfortable at all?
- Which project are you proud of?
 - What is {clean,ugly} code for you?
 - What are your expectations of today?
- What technical knowledge would you like to share?
- What is your experience with Docker so far?
- What is your goal for today's workshop?
- What is puzzling you about Docker?

About me

- `$> whoami` : DI Spiess-Knafl Peter
- Former Software Architect @ [bitmovin](#) and [skiline](#)
- Former Debian Maintainer (2014 - 2018)
- Founder of
 - technikrabe OG (2016 - 2019)
 - Spiess-Knafl Peter IT Solutions (since 2019) - [spiessknafl.at/peter](#)
 - hardcode GmbH (since 2023) - [hardcode.at](#)
 - [nobloat.org](#)
- GitHub [cinemast](#)
 - [libjson-rpc-cpp](#)
 - [json-rpc-cxx](#)

What is Docker?



docker

- A toolbox (UI) for managing containers
 - Build - creating OCI compliant Container images
 - Distribute - push images to registries
 - Deploy - pull images from registries
 - Run - create containers from images
- specifically **CGroups and Kernel Namespaces**

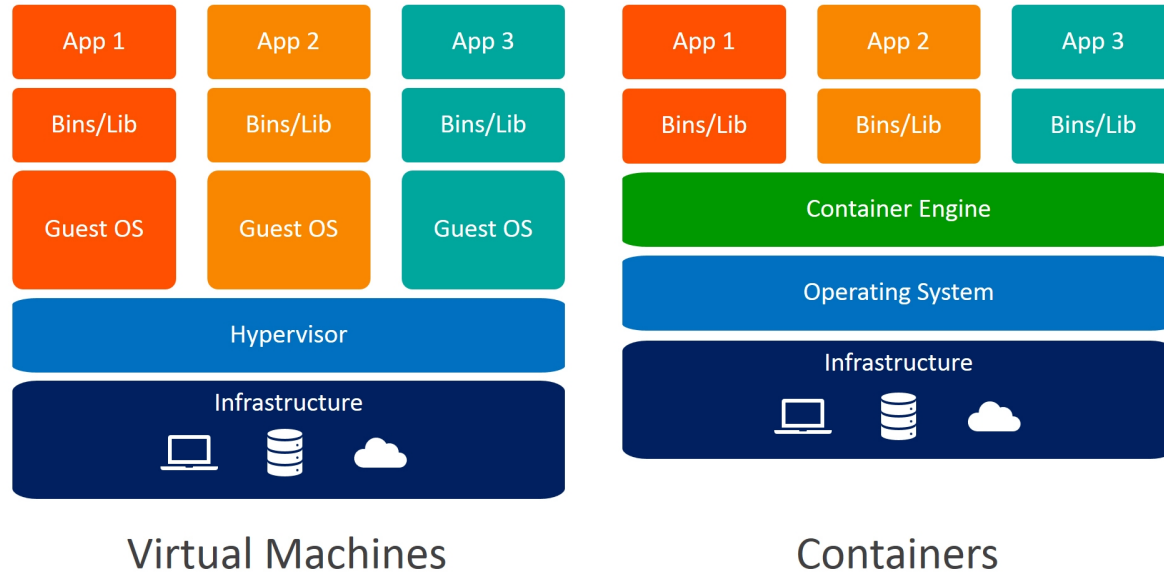
Why Docker?

- [Infrastructure as Code](#)
- Immutability and reproducibility
- Simple local development setups
- When there is no other easy mechanism to install/deploy software
- Has become the de-facto standard for container based environments
 - There is also [podman](#)
 - Both adhere [OCI \(Open Container Initiative\)](#)

What is a Container?

- CGroups/Namespace allows applications to run in **isolation** of each other, regarding syscalls.
 - Processes (PIDs)
 - Network (Ports)
 - Filesystem (Mounts)
- Sounds familiar?
 - Virtual Machines (VM) have a similar purpose

Containers vs. Virtual Machines



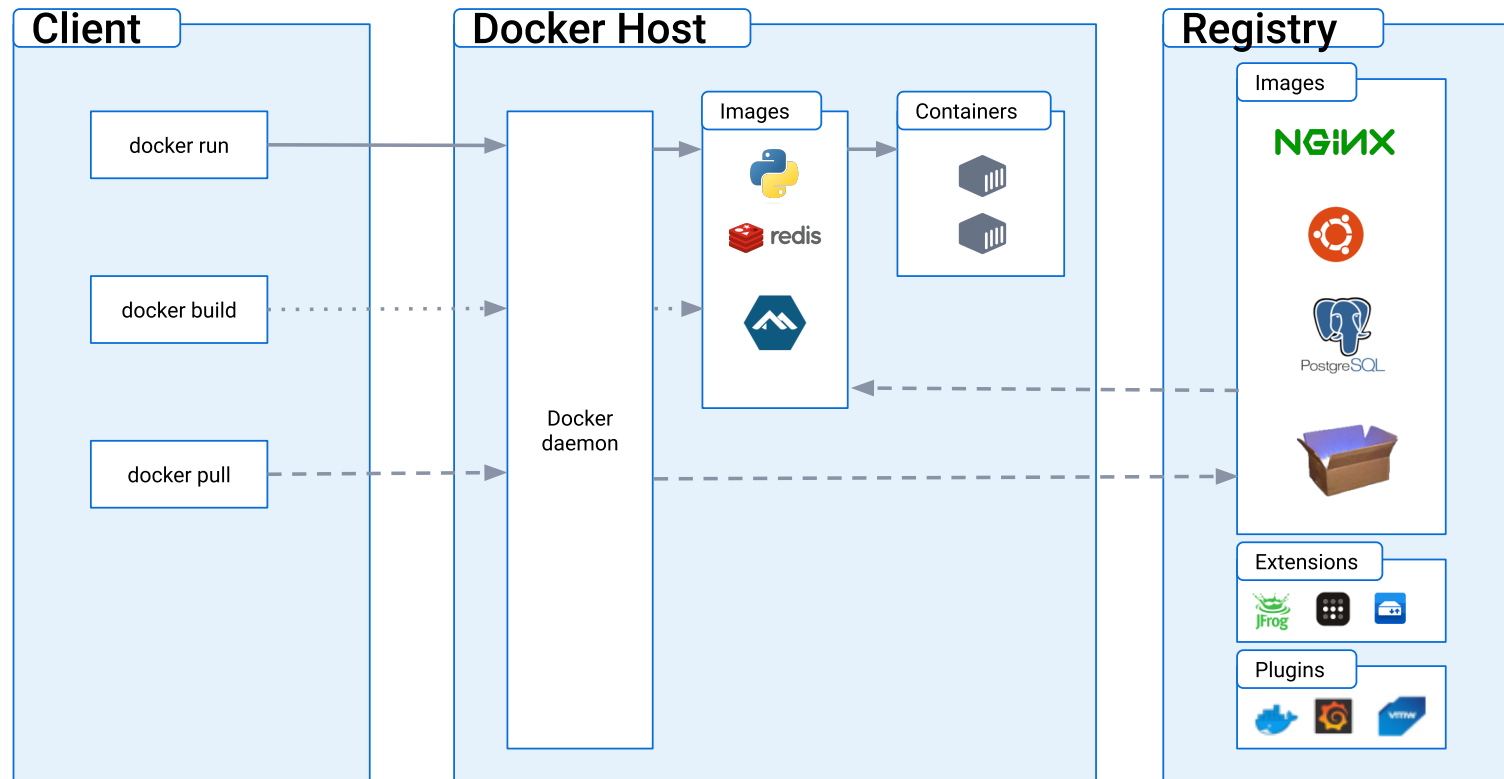
- All containers share the same Kernel
 - ... and the same **kernel vulnerabilities**
 - VMs provide a better isolation and security layer
- Abstraction at `syscall` level vs. Hardware level (VM)
 - `syscall` is the interface between the Kernel and your applications.

Key differences Containers/VMs

	Docker	VM
Resource usage	process overhead	full operating system
Scalability	Easy	Harder
Boot time	Sub (milli-)seconds	Minutes
Image size	Size of application + libs	Size of application + libs + full OS
Isolation	Kernel	Hypervisor

Docker concepts

- Docker Daemon: Core component, orchestrating all moving parts



Docker concepts

- **Containers**: Isolated process with its own
 - namespaces (network, filesystem, processes, etc.)
 - are created/instantiated out of *Images*
 - `docker run -it <image>`
- **Images**: Are basically a prepared "root" filesystem to "boot" a container from
 - Layered filesystem (each diff results in a new layer)
 - Images can inherit from base images (e.g. `FROM nginx`)
 - Tags: Images can be versioned using tags (e.g. `nginx:latest`)
 - Can be `docker build -t <repo>/<image>:<tag> .` using `Dockerfile`

Docker concepts

- **Registries**: Hold different Images
 - `docker pull <registry>/<image>:<tag>`
 - `docker push <registry>/<image>:<tag>`
- **Volumes**: Holds data should be persisted beyond a container's lifetime.
 - Can be `mounted` into containers
 - Named: `docker volume create myvolume1`
 - Bind: `docker run -v <host-directory>:<container-directory> -it <image>`
- **Capabilities**: Permissions a container can have
 - e.g. `SYS_ADMIN`, `NET_RAW`, `SYS_BOOT`

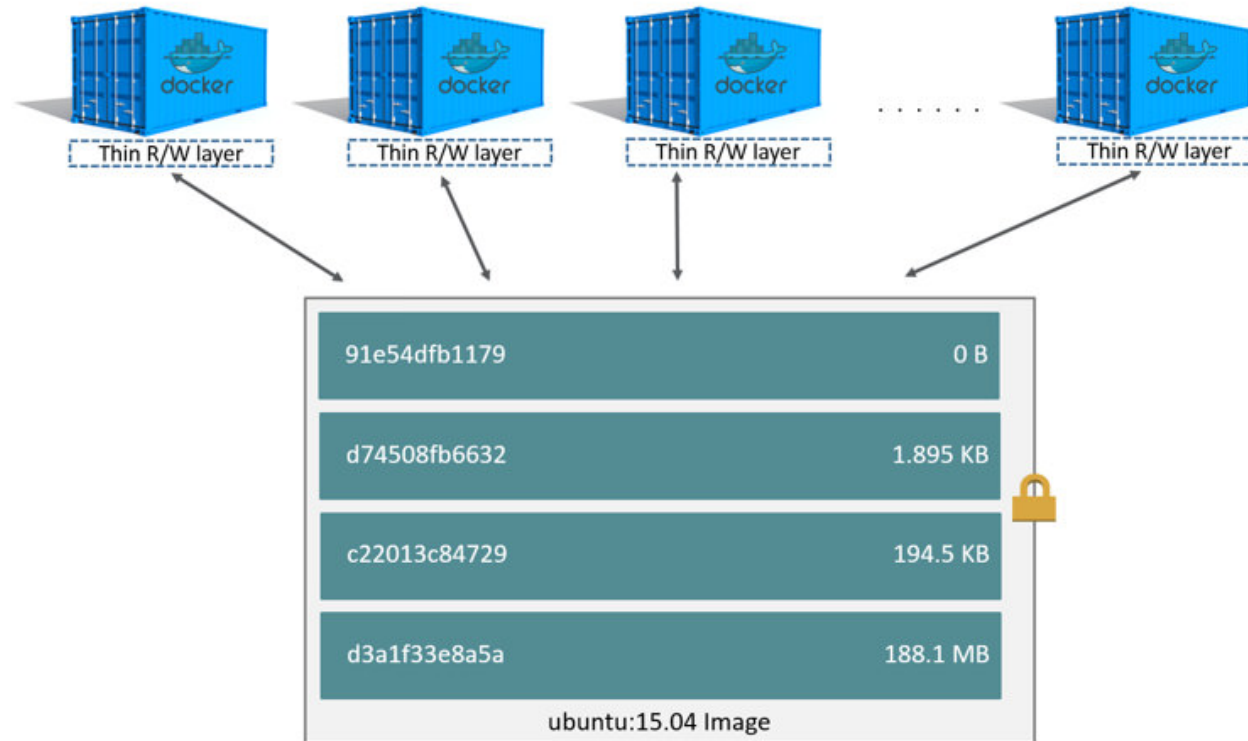
Building Images

```
docker build -t ghcr.io/my-org/my-service:3.8-alpine .
```

- `ghcr.io` - name/url of the registry to be pushed to, defaults to `hub.docker.io`
- `my-org` - project name
- `my-service` - image name
- `3.8-alpine` - image tag (version and variant), defaults to `latest`

Docker Image

- Consists of stacked R/O filesystem layers
- [Dockerfile](#): Build instructions to construct the image
- Dive example!



Dockerfile

- `FROM` defines the base image, first instruction of every `Dockerfile`
 - `FROM scratch` starts off with a completely empty image
 - `:alpine` and `:slim`
- `ADD` / `COPY` adds files to the image ([Copy vs. Add](#))
- `RUN` runs a command within the so far created image
- `ARG` build time variables
- `ENV` runtime environment variables (documentation)
- `EXPOSE` available ports that can be mapped (documentation)
- `VOLUME` defines volumes to be mounted

Dockerfile

- `USER` sets the user to be used in all subsequent layers. Default `root`
 - Containers should not run as `root` unless absolutely necessary.
- `WORKDIR` default directory the container will start in
- `HEALTHCHECK` execute the command as health-check mechanism
- `CMD` defines the command to be executed when the container is started
- `ENTRYPOINT` shell script the runs before the `CMD`
 - `/entrypoint.sh` must take `$1` as `CMD` parameter

Dockerfile best practices

- [Best practices](#)
- Start only one process per container
- Avoid custom `/entrypoint.sh` if not necessary
- Install only what you need during **runtime**.
 - if build- and runtime differ, use [multi-stage](#) builds
- `RUN rm /package.deb` will not reduce the file size -> layers
- Often changed parts should be last in the `Dockerfile`
 - layers are cached and only updated if necessary
 - each changed layers updates all above layers
- Mind the build context `.dockerignore`

Dockerfile security considerations

- Install only what is necessary
- Update base images before building `docker build --pull`
- [Security Best Practices](#)
- `docker scout cves <image>`
- `dive` - inspect docker layers interactively

Running containers

- Starting container with default CMD

- `docker run -it -e ENVIRONMENT=PROD -v /mnt/data:/data -p 8080:80 <image>`

- Run a command within the container (once)

- `docker run -it -e ENVIRONMENT=PROD -v /mnt/data:/data -p 8080:80 <image> <cmd>`

- Run a command within the container (once) and delete the container afterwards

- `docker run --rm -it -e ENVIRONMENT=PROD -v /mnt/data:/data -p 8080:80 <image> <cmd>`

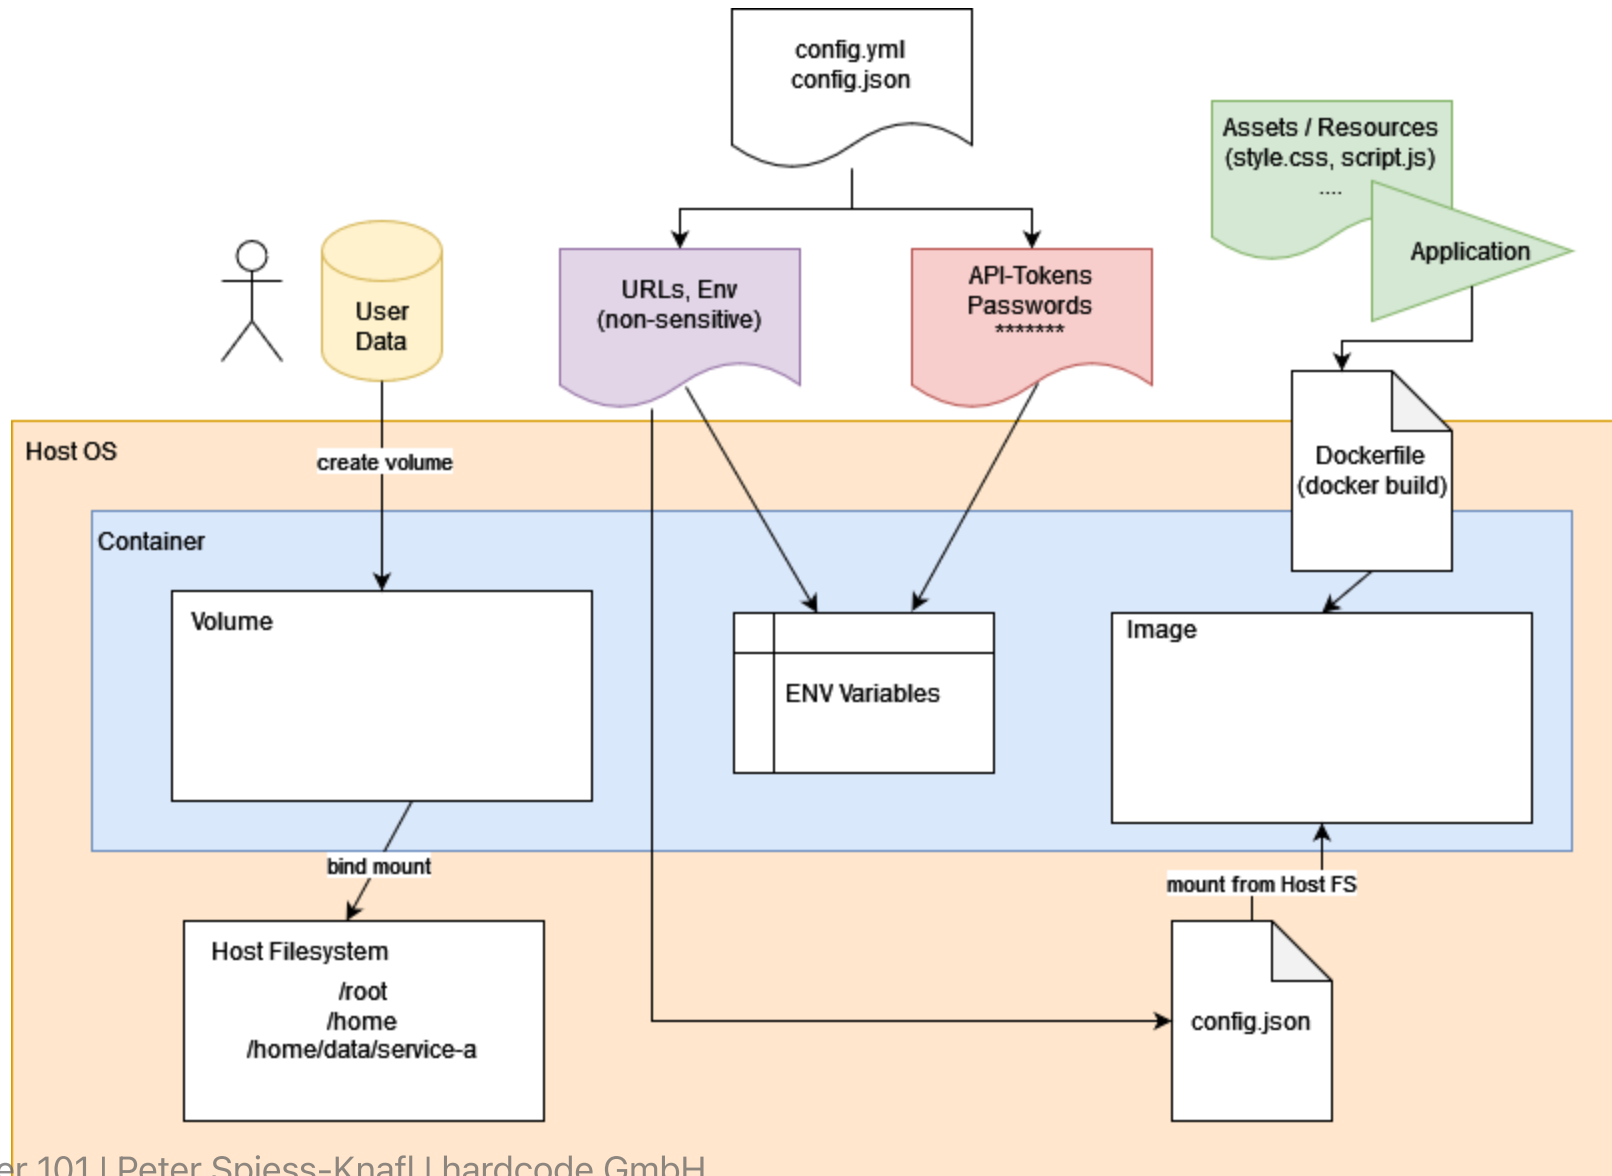
Running containers (2)

- Remembering and executing `docker run` with lots of parameters is cumbersome
- `docker compose` to the rescue, perfect for setting up local dev environments
 - Orchestrate multiple containers in a `compose.yml` file
 - Manage in one place
 - Containers, Environment variables, Port Mappings
 - Networks
 - Volumes
- Debugging:
 - `docker exec -it <container-id> /bin/sh`
 - `docker logs -f <container-id>`
- Cleanup: `docker system prune`

Where to put what kind of data?

- Usually different categories of data considering containers
 - Configuration (environment, API endpoint URLs)
 - Secrets (API Keys, database passwords, admin passwords, SMTP)
 - Application (bundled resources, source code, binaries)
 - User data (configuration, database rows)
- Where should they be kept?
 - What should be kept in the image?
 - What should be kept in the environment?
 - What should be kept in volumes?

Where to put what kind of data? (2)



Tools

- [Podman](#) - Alternative OCI compliant Docker alternative
- [dive](#) - Inspect docker images interactively
- [Docker Desktop](#) - Developer UI
- [Portainer](#) - Web interface for docker daemons
- [quay.io](#) - Registry alternative to Docker Hub
- [netshoot](#) - Network debugging tools in one image

docker compose

- Manage multiple containers at once
- Containers are defined in `compose.yml` or `docker-compose.yml`
- Very handy for local dev environments
 - 1 container that runs your PHP Code
 - 1 container that runs a MySQL database locally
- Manage containers, volumes, networks, ports, envs out of one file
- `docker compose up <service>`, `docker compose down`, `docker compose logs -f`
- `docker compose ps`, `docker compose pull`, `docker compose restart`

Cheat sheet

Resources

- [Docker documentation hub](#)
- [Docker cheat-sheet](#)
- [Dockerfile best practices](#)
- [Dockerfile reference](#)
- [Docker compose reference](#)
- [Docker curriculum](#)
- [Youtube tutorial](#)
- [Youtube crash course 1 hour](#)
- [The 12 factor App](#)

When not to use Docker?

- If you already have an easy to install software
 - e.g. `.deb` , `.rpm` packages
- If you already have single artifact that you can drop on a server
 - e.g. `.jar` files and Java already solved this problem
 - e.g. `go` binaries are statically linked and can just be run
- Automatic restart can also be solved using `systemd` units
- Adding another layer of abstraction won't make your life easier
 - You need at least a registry
 - You probably need a build pipeline
 - You need a deployment mechanism

Deploying PHP applications with docker

- Multiple scenarios possible
 - FROM php:8.2-apache
 - considered legacy
 - weaker performance compared to fpm
 - FROM php:8.2-fpm + nginx in one container
 - violates the one process per container rule, but definitely simplifies setup
 - FROM php:8.2-fpm + FROM:nginx + shared volume
 - complex setup but gives the greatest flexibility
 - suitable for "micro service" setups
 - nginx only runs once for all services
- Using Docker container as local environment for db + app

Deploying Node applications with docker

- Depends on your setup
- Frontend projects -> try to export as static sites and deploy into `nginx` or `s3` with cache
- Backend projects
 - Make sure you use multi-stage builds
 - Especially with typescript
 - Choose between `node`, `deno` and `bun`
 - `pm2` - 1 process per container rule, most of the time `pm2` can be replaced by native docker images without it.
- Using Docker container as local environment for db + app

Conclusion and Recap

- What is a Container and how does it differ from VMs?
- What is a docker image and how can I build my own?
- What is a registry?
- How can I create containers?
- How can Dockerfiles be optimized and why are layers important?
- Where should persistent data be kept?
- What useful tools are there for managing docker containers?

Thank you!